

A Formal Test of Assortative Matching in the Labor Market

John M. Abowd
Francis Kramarz
Sébastien Pérez-Duarte
Ian M. Schmutte

March 19, 2012

Any opinions and conclusions expressed herein are those of the authors and do not necessarily represent the views of the U.S. Census Bureau. All results have been reviewed to ensure that no confidential information is disclosed. This research uses data from the Census Bureau's Longitudinal Employer Household Dynamics Program, which was partially supported by the following National Science Foundation Grants SES-9978093, SES-0339191 and ITR-0427889; National Institute on Aging Grant AG018854; and grants from the Alfred P. Sloan Foundation. Abowd also acknowledges direct support from NSF Grants SES-0339191, CNS-0627680, SES-0922005, TC-1012593, and SES-1131848. The data used in this paper were derived from confidential data produced by the LEHD Program at the U.S. Census Bureau. All estimation was performed on public-use versions of these data available directly from the authors.

Abowd: Department of Economics, Cornell University, CES - U.S. Census Bureau, and NBER. Email: John.Abowd@cornell.edu

Kramarz: CREST(ENSAE). Email: francis.kramarz@ensae.fr

Pérez-Duarte: European Central Bank. Email: Sebastien.Perez-Duarte@ecb.europa.eu

Schmutte: Department of Economics, Terry College of Business, University of Georgia. Email: schmutte@uga.edu.

Abstract

A well-functioning labor market should sort skilled workers into jobs where skill is most valued. Employer-employee data reveal little sorting of workers to firms on unobservable correlates of earnings. We make three contributions toward explaining the apparent lack of sorting. First, we develop a general method for estimating matching models in employer-employee data. Second, we apply our method to a structural assignment model with coordination frictions (Shimer 2005). Assortative matching is present, but masked by selection of jobs from the population of job offers. Finally, we make the moments required for estimation, which are derived from restricted-access data, publicly available.

JEL Codes: J31, J24, E24.

Keywords: Wage Differentials; Human Capital; Skills;
Job Matching; Simulation Methods

1 Introduction

In this paper, we reconcile a tension between theoretical and empirical work on labor market sorting. Many models predict that more skilled workers should sort into more productive firms. However, analyses of matched employer-employee datasets find little evidence of positive assortative matching, as measured by the correlation between unobserved worker- and firm-specific components of earnings. The tension arises not because these findings are contradictory, but because they emerge from different approaches to addressing a similar problem. Our paper pulls together the theoretical and empirical literature by estimating a structural model of job assignment that makes specific predictions about the statistical wage decomposition.

The standard approach to measuring sorting in matched employer-employee data is to compute a statistical decomposition of wage rates into portions due to observable characteristics, unmeasured individual heterogeneity, and unmeasured employer heterogeneity. We follow the literature in referring to this wage decomposition, introduced to the economics literature in Abowd et al. (1999), as the ‘AKM decomposition’ throughout the paper. The AKM decomposition has two consistent features across a number of different datasets covering. First, it attributes substantial variation to the unmeasured individual and employer components.¹ Second, the correlation between worker and employer heterogeneity is generally negligible. There is less agreement on the sign of the statistical correlation between these components of heterogeneity in the observed samples of workers. Some find a small negative correlation (Abowd et al. 1999). Others find some evidence of positive correlation (Woodcock 2008; Abowd et al. 2003; Iranzo et al. 2008).

The negligible correlation between worker and employer heterogeneity components is a stylized fact in search of explanation. At face value, this finding indicates that labor mar-

¹For example, see Abowd et al. (1999) for France, Abowd et al. (2003) for the United States, Mortensen (2003) for Denmark, and Iranzo et al. (2008) for Italy.

ket matching is essentially random. Random assignment is inconsistent with the optimal sorting of workers to firms that occurs in frictionless assignment models, with or without comparative advantage (Becker 1973; Sattinger 1993). However, it is consistent with models of undirected search in the presence of labor market frictions. Roy (1951) posited an optimal sorting of workers with heterogeneous abilities that were differentially productive in different occupations. Employer heterogeneity resulted from different occupation mixes across industries and firm sizes. Mortensen (2003) demonstrates that the within occupation decomposition into unmeasured individual and employer components is just as strong as the between occupation decomposition, which is not consistent with Roy’s model. Mortensen (2003) offers a collection of alternative explanations based on search and recruitment costs among employers and workers with heterogeneous productivity.

Our paper makes three contributions that enhance our understanding of matching in the labor market. First, we estimate a general equilibrium model of job assignment with coordination frictions due to Shimer (2005) in matched employer-employee data for the U.S. In the absence of a clear economic model, it is difficult to interpret various statistics from the AKM decomposition. The structure of Shimer’s model yields a precise interpretation of the reduced-form wage decomposition for the nature of job assignment. Shimer’s model yields a wage offer function that depends upon both the ability of the individual and the productivity of the employer. Furthermore, it predicts which jobs will be filled in equilibrium. These are the two features needed to construct theoretical moments of the joint distribution of individual and firm heterogeneity. The structural model allows us to assess the extent to which the correlations we observe between individual and firm heterogeneity components are due to assortative matching, mismatch, and selection effects. Consistent with Shimer’s conjectures, we find that the small correlation between person- and firm- effects in the AKM decomposition understates the true extent of assortative matching.

Our second contribution is methodological. To estimate Shimer’s model, we demonstrate

that there is a direct mapping between the equilibrium quantities in his model and the estimated person and firm effects in the AKM decomposition. Estimation is then carried out using simulated method of moments. Our approach is applicable to a more general class of models with heterogeneous workers and firms. To use our approach, the model must yield (1) a wage offer function that depends only on the type of worker and the type of firm; (2) a prediction on set of employment relationships that are realized.

Our third contribution is to provide the complete set of moment matrices used in estimation. This contribution goes beyond standard best-practices regarding replicability because the underlying microdata used in this paper are confidential, and while access is open, in practice obtaining it can be quite time-consuming.² Researchers interested in labor market assignment can use these data in combination with our empirical strategy described above to estimate alternative models of assignment and matching. Our released data characterize the moments of wage components in the AKM decomposition, for twenty industrial sectors, stratified by employer size. There is no other publicly-available source of data based on matched employer-employee data with this degree of detail.³

We are not the first to address these questions. Lentz (2010) and le Maire and Scheur (2008) consider frictional assignment models with endogenous search intensity. Both find that the AKM decomposition can understate the true extent of sorting, depending on the underlying specification, but these findings are based on simulations rather than direct estimation. le Maire and Scheur (2008) and Woodcock (2008) estimate partial equilibrium models that include match effect and show that there may be some negative bias in the correlation between person- and firm- effects in the AKM decomposition when the match effect is omitted. Mendes et al. (2010) find evidence of positive assortative matching by directly

²Our data are from the LEHD program of the U.S. Census Bureau, which we describe in full detail below. To learn more about how to access the confidential microdata, see <http://www.census.gov/ces/index.html>

³Our data and code are freely available online for download from http://schmutte.myweb.uga.edu/content/AssortMat_DataArchive_2012March7.zip.

estimating firm-specific production functions, and comparing firm productivity to worker effects from the AKM decomposition. Lopes de Melo (2009) has the only paper, aside from this one, that estimates a general equilibrium model using matched employer- employee data with the objective of explaining the small correlations observed in the AKM decomposition. Eeckhout and Kircher (2011) discuss the general problem of identifying sorting in wage data when wages are not monotonic in the type of the firm, which can generically arise in many models, including our coordination friction model.

We estimate the structural model in each of twenty major NAICS industrial sectors. For exposition, in the paper we present results from two representative sectors: Manufacturing and Professional Services. These two sectors characterize the equilibria we observe most frequently.⁴ In the Manufacturing sector, we find that most workers are of low ability. Among potential jobs, only a small fraction are of high productivity. The production technology is such that it does not matter who takes which job – a fact that favors negative assortative matching. In equilibrium, low-productivity worker queue for jobs in the high-productivity sector, but are guaranteed work in the low-productivity sector. They receive compensation from the high-productivity firms to accept unemployment risk associated with queuing for those jobs. The negative correlation between the offered wages in the two sectors is offset by the selection of matches.

In the Professional Services sector, we estimate that most workers are of high ability, and that technology favors positive assortative matching. Again, only a very small fraction of available vacancies are of high productivity. Employers with high productivity jobs offer very high wages to low ability workers as compensation against the risk of being employed because they are not selected from the queue. Thus, in the distribution of wage offers, there is a strong negative correlation. However, in the distribution of realized wages, there is a weak positive correlation because most of the realized matches involve high ability workers,

⁴The complete set of results across all twenty sectors are available from the authors upon request.

with almost all low ability workers in low productivity jobs.

In both cases, the model implies that observed small correlations actually mask much stronger assortative matching. Moreover, the model indicates that the problem is one of sample selection. We only observe the matches that are realized, rather than a random sample of matches from the full distribution of potential offers. Finally, we observe that, although it was not a feature of the model, our findings are related to dual labor market theory. Workers can either find immediate employment in a large low-productivity sector, or queue at some risk for more desirable employment in a thin high-productivity sector.

Section 2 describes the features of Shimer’s assignment model with coordination frictions that are essential for our analysis. Section 3 presents the input data from the LEHD program and describe how it is processed prior to estimation of the structural model. Section 4 summarizes the estimation procedure, and Section 5 the results of the estimation on LEHD data for the Professional, Scientific and Technical Services and Manufacturing sectors. Section 6 concludes.

2 Theoretical Assignment Model

In this section we introduce notation and present the key features of Shimer’s coordination friction model that are required for our analysis. Shimer (2005) has many details and implications that we omit.

2.1 Job Assignment with Coordination Frictions

The economy consists of workers and employers. Each employer has one vacancy to fill by hiring one worker. Employers choose wage offers for each worker. Workers then choose where they want to apply for a job. Based on the pool of applicants, the employers hire one worker to fill the vacancy. Workers and employers only differ in productive type. There are M types

of worker and N types of employers. When a worker of type m is hired into a job of type n , the output of production is $x_{m,n}$.

A decentralized equilibrium of the coordination friction economy is characterized by application strategies for all workers and wage offers from all employers to each type of worker. It turns out that this equilibrium depends only on the nature of the production technology and the number of each type of worker and of each type of employer. Out of a total measure μ of workers in the economy, we assume that μ_m of them have productive type $m \in \{1 \dots M\}$. Likewise, out of a total measure ν of employers in the economy, ν_n of them have productive type $n \in \{1 \dots N\}$. Although it will not matter in the final equilibrium, we follow Shimer in naming individual workers by (m, i) and individual employers by (n, j) where the first term is the type of the agent, and the second term is the agent's name. Thus, each worker or employer can be uniquely identified.

2.2 Equilibrium Wage Offers, Applications and Hiring

2.2.1 The Worker's Problem

Suppose each employer, (n, j) , has chosen its menu of wage offers. Employers cannot offer different wages to workers of the same type. Consequently, a worker of type m must choose an application strategy based on the set of offers, $w_{m,(n,j)}$, across all (n, j) . The model assumes that a worker can make only one application, so a feasible strategy for worker (m, i) is a probability distribution $p_{m,i}$ that specifies the probability with which she applies to each employer.

Since all workers of type m follow the same strategy in the coordination friction economy, and since we assume that the number of workers of each type is very large, it is easier to think in terms of the queue of applicants of type m to each employer. That is, the number of applications from type m workers to any employer is approximately Poisson with parameter

$$q_{m,(n,j)} = p_{m,(n,j)}\mu_m.$$

The equilibrium queue lengths must leave all workers of type m indifferent between applying for any two jobs. To satisfy this condition for each type of worker, all jobs for which there is a positive probability of application have the same expected income. In addition, all jobs for which the worker of type m does not apply have strictly less expected income. The expected income, $y_{m,(n,j)}$, from any job for which $q_{m,(n,j)} > 0$ is:

$$y_{m,(n,j)} = e^{-Q_{m+1,(n,j)}} \frac{1 - e^{-q_{m,(n,j)}}}{q_{m,(n,j)}} w_{m,(n,j)} \quad (1)$$

where $Q_{m+1,(n,j)} = \sum_{m'=m+1}^M q_{m',(n,j)}$ is the measure of workers with productive type at least $m + 1$. This is the expected number of applications from workers who are strictly better than a worker of type m . Using the properties of the Poisson distribution, the expected income from (n, j) is the probability of getting the job times the wage, where the probability of getting the job is the probability that there are no applications from a better worker (the first term on the right-hand side of (1) times the probability that the worker actually ends up applying for the job (the term in the numerator) and is chosen from among all other applicants of the same type (the term in the denominator). As long as there are expected incomes, $y_{m,n}$, and queue lengths that satisfy (1), workers have no incentive to deviate.

2.2.2 The Employer's Problem

Employers choose wage offers for workers of each productive type and a hiring protocol that maximizes expected profits given the queueing behavior of workers. Shimer (2005) shows that the optimal hiring strategy is to choose the applicant of highest productive type. When pursuing the optimal strategy, the maximized expected profit of employer (n, j) is given by

$$\sum_{m=1}^M e^{-Q_{m+1,(n,j)}} [1 - e^{-q_{m,(n,j)}}] [x_{m,n} - w_{m,(n,j)}]. \quad (2)$$

The employer's optimal wage offers maximize the difference between each worker type's productivity and wage rate times the probability that a worker of that type will apply.

2.2.3 Competitive Search Equilibrium

Shimer defines the equilibrium for this economy as follows:

Definition 1 A *Competitive Search Equilibrium* consists of wage offers, w , queue lengths, q , and expected incomes, y , chosen so that employer's profits (2) are maximized, worker's expected incomes are given by (1), and the expected number of applications from type m workers does not exceed the total measure of such workers, μ_m :

$$\mu_m = \sum_{n=1}^N \int_0^{\nu_n} q_{m,(n,j)} dj. \quad (3)$$

The competitive equilibrium turns out to be unique. Furthermore, the equilibrium queue lengths, $q_{m,n}$, are the same as those that would be chosen by a social planner who chooses queue lengths, or equivalently, application probabilities for workers. The social planner is restricted to queue lengths that do not condition on individual identities. The goal of the social planner is to maximize expected aggregate output. In this setting, it turns out that the optimal queue lengths will satisfy the first-order conditions:

$$\lambda_{m,n} \geq e^{-Q_{(m,n)}} x_{m,n} - \sum_{m'=1}^{m-1} e^{-Q_{m'+1,n}} (1 - e^{-q_{m',n}}) x_{m',n} \quad (4)$$

where $\lambda_{m,n}$ is the expected employment of type m workers in type n employers. The dependence on the individual's identity, i , and the employer's identity, j , have been suppressed because they are not relevant. Given these queue lengths, and the requirement that workers'

expected incomes satisfy (1), the equilibrium wages must be:

$$w_{m,n} = \frac{q_{m,n}e^{-q_{m,n}}}{1 - e^{-q_{m,n}}} \left[x_{m,n} - \sum_{m'=1}^{m-1} e^{-Q_{m'+1,n}} (1 - e^{-q_{m',n}}) x_{m',n} \right]. \quad (5)$$

Equations (4) and (5) provide all the information that is required for our estimation strategy. Given assumptions about the production technology, $x_{m,n}$, and the measure of each type of employer and worker, wages and queue lengths are determined. Although the equilibrium queue lengths have no observable counterpart, we can use them to derive the expected number of workers of type m that ultimately become employed by employers of type n :

$$\lambda_{m,n} = \nu_n e^{-Q_{m+1,n}} (1 - e^{-q_{m,n}}). \quad (6)$$

2.3 Empirical Predictions and Sorting

In equilibrium, as in the frictionless model, the nature of job assignment depends on the properties of the production function. Hence, the model is still informative about the nature of assortative matching in the economy. Unlike the frictionless model, however, the coordination friction equilibrium has several properties that are more consistent with the stylized facts characterizing labor markets. These equilibrium properties are also enough to generate a structural connection between the underlying model primitives and the estimated person and employer effects from a linear decomposition of log earnings. We exploit this structural link to estimate the structure of the model.

With respect to the nature of assortative matching, Shimer derives several strong results. First, if the production function is supermodular, Shimer's Proposition 3 shows that $Q_{m,n}$ is strictly increasing in n when it is positive. In addition, a more productive job is more likely to be filled, a worker is less likely to obtain a more productive job conditional on applying for it, and a worker's wage is increasing in the employer's productivity. This result holds

for production technologies that exhibit comparative advantage for higher-ability workers in higher-productivity jobs. It also holds, as is well known, for a production technology that exhibits no comparative advantage. Specifically, one in which output is multiplicative in the ability and productive types of the worker and job. A converse result holds when there is a comparative advantage for low ability workers in high-productivity jobs. In that case, $Q_{m,n}$ is decreasing in n so that the assortative matching again corresponds to comparative advantage. We estimate the structural model assuming either that the production technology exhibits no comparative advantage or that the technology exhibits comparative advantage for low ability workers in high productivity jobs. These specifications yield the minimal differentiation sufficient to capture model uncertainty as to whether matching is positively or negatively assortative.

On their own, these results do not distinguish the coordination friction economy, since the frictionless model also predicts that job assignment will be in accord with comparative advantage in production. However, the coordination friction model makes more reasonable predictions about equilibrium unemployment and vacancies, and predicts a certain amount of mismatch, which eliminates the untenable “perfect correlation” prediction for worker and employer types from comparative advantage models without the coordination friction. Since workers randomly choose the employer to which they will apply, there is, in general, a positive probability that multiple workers apply for the same job—leaving all but one unemployed. There is also a positive probability that no one will apply for any particular job, leaving it vacant. Therefore, there will be simultaneous unemployment and job vacancy for workers and jobs of all types. Furthermore, even among employed workers and occupied vacancies, there will be mismatch in the sense that the particular set of matches realized in the economy could be improved upon in the sense of maximizing total output, or in terms of aligning assignments with comparative advantage.

Finally, the coordination friction equilibrium has two properties that are essential to the

estimation strategy. In equilibrium, the wage offer function violates the law of one price. Different jobs pay identically able workers different wages. That is, there is a structural employer effect in the wage. Shimer shows that the wage offer will generally be increasing in the productive type of the employer. The equivalent result does not hold for workers, however. It is possible for an employer to offer workers with low ability higher equilibrium wage than workers with higher ability. This establishes that it is possible to observe a negative correlation between person and employer effects from the empirical earnings decomposition even when there is positive assortative matching on the unobserved productivity types.

Note, however, that the model does not generally imply that log earnings can be additively decomposed into a person and employer effect. But this implication is not required by our estimation procedure. Our approach is motivated by the observation that the person and firm effects from the statistical decomposition of log earnings in a linear model will be complicated transformations of the underlying productivities and abilities along with factors determining the supply of and demand for worker skill. As we show below, Shimer’s coordination friction model provides explicit formulas for these transformations that we can identify through restrictions on observable moments of the the joint distribution of the estimated person and firm effects from the earnings decomposition.

3 Data

Our estimation procedure has two stages. First, we use matched employer-employee data from the U.S. Census Bureau’s LEHD Program to estimate the parameters of a statistical decomposition of log earnings into components due to individual heterogeneity, employer heterogeneity, and time-varying observables. Next, we estimate the structure of the assignment model from its predictions about the joint distribution of the person and employer effects. In the structural estimation, we use the distribution of person and employer effects stratified

by deciles of the firm size distribution. This provides an additional source of variation to help identify the model in the absence of firm-level data on productivity. We estimate the model separately for the Manufacturing (NAICS 31-33) and the Professional, Scientific, and Technical Services (NAICS 541) sectors.

3.1 The Statistical Decomposition

Following Abowd et al. (1999), we decompose the annualized wage in each job as:

$$\log w_{i,t} = x_{i,t}\beta + \theta_i + \psi_{J(i,t)} + \varepsilon_{i,t} \quad (7)$$

where

- $J(i, t)$ is the dominant employer of individual i at time t ;
- θ_i is the individual effect;
- $x_{i,t}\beta$ is the effect of individual labor market experience, tenure, and labor market attachment, interacted with sex; and
- $\psi_{J(i,t)}$ is the firm effect.

Abowd et al. (2002) document the identification strategy and the details of implementing the full least squares solution for estimating (7).

We estimate (7) on matched employer-employee data from the LEHD Program of the U.S. Census Bureau. The LEHD program uses administrative data from state Unemployment Insurance wage records and ES-202/QCEW establishment reports that cover approximately 98 percent of all non-farm U.S. employment. Our estimation sample is based on a snapshot of the LEHD data infrastructure that includes data from twenty-eight states collected between 1990 and 2004. A detailed description of the input data is found in Appendix B.

3.2 The Data Used in the Structural Estimation

The estimates of the assignment model are based on data, D , of the form:

$$(\hat{\theta}_i, \hat{\psi}_j, s_\ell)_\ell$$

where $\hat{\theta}$ and $\hat{\psi}$ are estimates of the true person and firm effects from (7); i and j are the identities of the worker and firm contributing to the ℓ th observation; and $s_\ell = (s_1, \dots, s_{10})$ is an indicator for the decile of the firm size distribution into which the employer falls. We assume that $\ell \in \{1, \dots, N\}$ and define $\sum_{\ell=1}^N s_\ell = (N_1, \dots, N_{10})$.

We compute the empirical first and second moments of the joint distribution of θ and ψ within employer size decile for each of the two NAICS sectors. The unit of analysis is the individual job history observation in the matched data. The deciles of the employer size distribution are computed for every year of the data, and employers are assigned to the decile that they occupy for each year. We use a bootstrap procedure to compute the covariance matrix for the estimated first and second moments. We resample the data from each sector 1,000 times and compute the first and second moments of the joint distribution of $\hat{\theta}$ and $\hat{\psi}$ in each sample. The estimated covariance matrix is the covariance matrix of these 1,000 estimates.

One of the contributions of this paper is that we are able to make the data used in estimation publicly available, even though they are based on confidential microdata.⁵ It is important to note the measures taken to protect the confidentiality of respondents providing the underlying micro-data. The U.S. data are protected by a distribution-preserving noise-infusion procedure. Each element of the microdata has been distorted in such a way that the individual observations are protected, yet the aggregate statistics are still valid for our

⁵The moment matrices described in the preceding paragraphs are available upon request from the authors and online at http://schmutte.myweb.uga.edu/content/AssortMat_DataArchive_2012March7.zip.

proposed analysis. For complete details of the noise infusion procedure and the way in which it preserves analytical validity, see Abowd et al. (2012).

We have estimated our model on all twenty major NAICS sectors. For exposition, in the remainder of the paper, we focus our discussion on two sectors: Manufacturing and Professional, Scientific, and Technical Services. These sectors are interesting in their own right, but also capture the two main profiles of estimated equilibria that we observe across all sectors. Detailed results for all sectors are available in an online Web Appendix,⁶ and upon request from the authors.

Figure 1 provides a visual summary of data from the Manufacturing and Professional, Scientific and Technical Services sectors. For Professional Services, Figure 1a shows that employment in this sector is concentrated in smaller firms. Figures 1c and 1e show two different kinds of evidence about the relationship between the estimated person and firm effects. The first figure plots the empirical average person and firm effects within each decile of the employer size distribution. The average firm effect is increasing in employer size, consistent with the widely documented presence of a firm size effect in the wage equation. Note that here the firm-size effect is preserved even controlling for the effects of personal heterogeneity. Abowd et al. (2003) note that the firm-size effect is indeed driven almost entirely by firm heterogeneity in wages with little of the effect caused by sorting of highly paid workers into larger firms. Here, we see that the average level of the person effect does not change much across deciles of the employer size distribution. The standard errors on these estimates are so small as to be invisible on the plot. Figure 1e shows the empirical correlation between $\hat{\theta}$ and $\hat{\psi}$ within each decile of firm size. The correlations are generally small in absolute value, ranging between about -0.2 and 0.25 . Our model will fit estimated correlations that try to capture the scale of these correlations as well as the variation across size classes. Here, we observe a hump-shaped relationship between employer size decile and

⁶http://schmutte.myweb.uga.edu/content/AssortMat_webAppendix_AllSectors.pdf

the correlation between $\hat{\theta}$ and $\hat{\psi}$. In terms of the model, this is feasible only if there is some considerable variation in the production technology and demand across the size classes.

Figures 1b, 1d and 1f show the corresponding picture for the Manufacturing sector. The distribution of jobs is more concentrated in firms of intermediate size. As in the Professional, Scientific and Technical Services sector, we see clear evidence of a firm-size effect on earnings. Here, though, there is a more puzzling relationship between firm size and the average person effect. The relationship is clearly non-linear, but its shape defies simple interpretation. Here, the correlations are smaller, tending close to zero, with the exception of a relatively strong negative correlation between $\hat{\theta}$ and $\hat{\psi}$ for the smallest firms.

4 Estimation Procedure

In this section we show that Shimer’s model generates restrictions on the joint distribution of the estimated person and employer effects, $\hat{\theta}$ and $\hat{\psi}$ from the decomposition of log earnings that are sufficient to identify the free parameters of the model. We analyze the conceptual log earnings decomposition on the matched employer-employee data that would be generated in the coordination friction economy. We derive formulas for the expectation of estimated person and firm effects in terms of equilibrium quantities of the model. Using these theoretical analogues to the estimated person and firm effects, we derive formulas for the expected value of the empirical second moments of $\hat{\theta}$ and $\hat{\psi}$. The full set of estimating equations uses these restrictions on the second moments of the joint distribution of $\hat{\theta}$ and $\hat{\psi}$ as well as restrictions on the distribution of realized matches across deciles of the firm size distribution. The section then goes on to describe how we parameterize the model and estimate it computationally.

This section contains one of the major contributions of this paper. The empirical strategy we derive is not specific to Shimer’s model. The key equations, Equations 8, 9, and 10, are

general, and relate the AKM decomposition to any model that makes predictions on wage offers and realized matching sets. The estimating equations that follow in Section 4.1.2 are equally general.

In bringing Shimer's coordination friction model to the data, we make two key assumptions, which we describe here, and discuss further at the end of the section. The first assumption is interpretive. The AKM decomposition is based on an analysis of dynamic panel data, but the theoretical model is static. Our analysis implicitly requires that we can think of the observed data as being generated by sampling repeatedly from the static equilibrium of a particular coordination friction economy. While this assumption is strong, we will argue at the end of this section that it is a defensible approach to structural estimation for this problem. Second, Shimer's model is based on bilateral matching, and is, like most models in this literature, silent about the agglomeration of jobs into firms. We assume that firm size is an observable correlate of employer type.

4.1 Method of Moments Estimation

4.1.1 Theoretical Wage Decomposition

We consider two types of workers, so $M = 2$. In the coordination friction economy, for any given employer all workers of the same type are equally productive and by assumption are equally paid. The wage offer functions for the case $M = 2$ are given by:

$$\begin{aligned} w_{1,n} &= \frac{q_{1,n}e^{-q_{1,n}}}{1 - e^{-q_{1,n}}}x_{1,n}, \\ w_{2,n} &= \frac{q_{2,n}e^{-q_{2,n}}}{1 - e^{-q_{2,n}}}(x_{2,n} - (1 - e^{-q_{1,n}})x_{1,n}) \end{aligned}$$

Recall that $\lambda_{m,n}$ is the expected number of m type workers matched to employers of type n and $w_{m,n}$ the corresponding wage. For $M = 2$ we have $\lambda_{1,n} = \nu_n e^{-q_{2,n}}(1 - e^{q_{1,n}})$ and

$$\lambda_{2,n} = \nu_n(1 - e^{q_{2,n}}).$$

To generate the theoretical analogues to the estimated person and firm effects, we perform a decomposition of log earnings into portions due to worker and employer type based on the equilibrium wage offers, w , and expected realized matches, λ , given model primitives (μ, ν, f) :

$$\log w_{m,n} = \theta_m + \psi_n$$

where dependence on i and j has been suppressed because it is irrelevant. There are only two worker ability effects to estimate and N employer productivity effects. We restricted our attention to models that produce equilibrium matches so that the decomposition is identified. The worker and firm effects are the solution to the weighted least squares problem

$$(\hat{\theta}_1, \hat{\theta}_2, \hat{\psi}_1, \dots, \hat{\psi}_N) = \arg \max \left[\sum_{m,n=1}^{2,N} \lambda_{m,n} (\log w_{m,n} - \theta_m - \psi_n)^2 \right]. \quad (8)$$

The worker and employer effects are identified up to a constant Λ (see Abowd et al. (1999) and Abowd et al. (2002)). Straightforward computations (see Appendix) yield

$$\hat{\theta}_i = \frac{\sum_n \frac{\lambda_{1n}\lambda_{2n}}{\lambda_{1n} + \lambda_{2n}} \log w_{i,n}}{\sum_n \frac{\lambda_{1n}\lambda_{2n}}{\lambda_{1n} + \lambda_{2n}}} + \Lambda, \quad i = 1, 2 \quad (9)$$

$$\hat{\psi}_n = \frac{\lambda_{1n}}{\lambda_{1n} + \lambda_{2n}} (\log w_{1,n} - \theta_1) + \frac{\lambda_{2n}}{\lambda_{1n} + \lambda_{2n}} (\log w_{2,n} - \theta_2) - \Lambda. \quad (10)$$

For consistency with the assumption used in the empirical estimation, we set $\Lambda = 0$.⁷

4.1.2 Estimating Equations

The estimating equations are based on the second moments of the $(\hat{\theta}, \hat{\psi})$ distribution resulting from solving (8), stratified by deciles of the firm size distribution s , and also from the expected

⁷See Appendix for the general case with a continuum of firms.

fraction of filled jobs within firms of each size class, that is the first moments of s . Within each size class, k , we form a vector of second central moments as:

$$\bar{m}_k^C(D) = \bar{m}^C(\hat{\theta}, \hat{\psi} | s_k = 1) = \frac{1}{N_k} \begin{bmatrix} (\hat{\theta} - \bar{\theta}_k)^T (\hat{\theta} - \bar{\theta}_k) \\ (\hat{\theta} - \bar{\theta}_k)^T (\hat{\psi} - \bar{\psi}_k) \\ (\hat{\psi} - \bar{\psi}_k)^T (\hat{\psi} - \bar{\psi}_k) \end{bmatrix}$$

where $\bar{\theta}_k$ and $\bar{\psi}_k$ are the within size-class means for decile k . In addition, the model makes predictions on the number of jobs observed in each size class:

$$\bar{m}^S(D) = \frac{1}{N} \sum_{\ell=1}^N s_{\ell}.$$

The 40 empirical moments are compactly expressed as

$$\bar{m}(D) = \begin{bmatrix} \bar{m}_1^C(D) \\ \vdots \\ \bar{m}_{10}^C(D) \\ \bar{m}^S(D) \end{bmatrix}.$$

The first 30 elements of $\bar{m}(D)$ are the second moments listed sequentially by size class and the last ten elements are the employment counts in each decile.⁸

We now show that the model structure is sufficient to specify a function that delivers moment restrictions corresponding to $\bar{m}(D)$ in terms of the underlying primitives of the model, (μ, ν, f) :

$$E(\bar{m}(D)) = g(\mu, \nu, f).$$

⁸For the overall economy, these counts are obviously equal in the real data; however, within each estimating sector (NAICS sector) they will not be equal.

We illustrate the manner in which the model determines the expectation of the empirical covariance between $\hat{\theta}$ and $\hat{\psi}$ within size-class k . The formulas for the variances are derived identically.

$$\begin{aligned}
E[cov_k(\hat{\theta}, \hat{\psi})] &= E \left[\frac{1}{N_k} \sum (\hat{\theta}_i \hat{\psi}_j) - \bar{\hat{\theta}}_k \bar{\hat{\psi}}_k \right] \\
&= \sum (\lambda_{11k} \theta_1 \psi_{(k,1)} + \lambda_{12k} \theta_1 \psi_{(k,2)} + \lambda_{21k} \theta_2 \psi_{(k,1)} + \lambda_{22k} \theta_2 \psi_{(k,2)}) \\
&\quad - \frac{(\lambda_{11k} + \lambda_{12k}) \theta_1 + (\lambda_{21k} + \lambda_{22k}) \theta_2}{(\lambda_{11k} + \lambda_{12k} + \lambda_{21k} + \lambda_{22k})} \frac{[(\lambda_{11k} + \lambda_{21k}) \psi_{(k,1)} + (\lambda_{12k} + \lambda_{22k}) \psi_{(k,2)}]}{(\lambda_{11k} + \lambda_{12k} + \lambda_{21k} + \lambda_{22k})}.
\end{aligned} \tag{11}$$

Using formulas (9)-(10) the covariance and variances are thus functions of the expected number of each type of worker per firm $\lambda_{m,n}$ and of the wage $w_{m,n}$. We derive the theoretically predicted number of filled jobs in each size class

$$g^S(\mu, \nu, f) = \begin{bmatrix} \frac{(\lambda_{1,(1,1)} + \lambda_{1,(1,2)} + \lambda_{2,(1,1)} + \lambda_{1,(1,2)})}{\sum_{s=1}^{10} (\lambda_{1,(s,1)} + \lambda_{1,(s,2)} + \lambda_{2,(s,1)} + \lambda_{1,(s,2)})} \\ \vdots \\ \frac{(\lambda_{1,(10,1)} + \lambda_{1,(10,2)} + \lambda_{2,(10,1)} + \lambda_{1,(10,2)})}{\sum_{s=1}^{10} (\lambda_{1,(s,1)} + \lambda_{1,(s,2)} + \lambda_{2,(s,1)} + \lambda_{1,(s,2)})} \end{bmatrix}.$$

These two variables are indeed determined in the Shimer economy.

4.2 Implementation

Our estimation procedure is based on matching the empirical second moments of the joint distribution of $\hat{\theta}$ and $\hat{\psi}$ to their theoretical counterparts. We obtain the optimal minimum distance estimator of the model parameters. Parametric restrictions on the model primitives aid in the identification of the model and also in easing the computational burden of estimation. Our inferences on the nature of assortative matching are based on estimates from several different specifications of the production technology. We use several representative technologies to capture the assumptions of comparative advantage for low-skilled workers in high productivity jobs (LCA) and no comparative advantage (NCA). We also fit the model

under a quasi-nested combination of these production technologies.

4.2.1 Parametrization

We have 30 empirical moments from the covariance matrix for each employer size decile and an additional 10 moments from the distribution of realized jobs across size classes. With the model restricted to two latent skill groups for workers and two latent productive types for firms within each size class, the most general specification of the model primitives would have 38 parameters associated with the production technology, 1 for the distribution of workers across skill groups, and 18 for the number of vacancies for jobs of each type. To satisfy the order condition for identification, we specify parametric restrictions that express assumptions about the relationship between firm size, vacancies, and productivity. Each of these is discussed in turn below.

There are $M = 2$ types of workers; the first parameter is α , the share of the type-1 persons in the population (α is between 0 and 1). The total number of employed workers J is known, but not the size of the labor force. We calibrate the unemployment rate at 6% (the results appear to be robust with respect to this parameter). Hence $\mu_1 = \alpha J / 0.94$, $\mu_2 = (1 - \alpha) J / 0.94$.

We consider 10 different size classes with two latent productive types per size class for a total of 20 different types of firms. The index n of the firm is (s, j) , where $s = 1, \dots, 10$ indicates the size class and $j = 1, 2$ indicates the type of the firm. The type of the firm represents a variation in the production function between productive and unproductive firms. The size categories are constructed so that each size category contains the same number of employed workers, J_s . The number of jobs (vacant and filled) is not observed, but the number

of employed workers by size category is observed. We model the number of jobs as

$$\begin{aligned}\nu_{s,j} &= \pi_{s,j} [1 + \exp(\gamma_0 + \gamma_1 \log(\sigma_s))] J_s \\ \pi_{s,2} &= 1 / [1 + \exp(\delta_0 + \delta_1 \log(\sigma_s))] \\ \pi_{s,1} &= 1 - \pi_{s,2}\end{aligned}$$

where $\nu_{s,j}$ is the number of jobs in firm (s, j) , $\pi_{s,j}$ is the share of type- j firms in the size category s , σ_s is the average size of the firms in size category s (σ_s is increasing in s).

The productive type of a job has two dimensions, the size-class, s , of the firm into which it is aggregated, and a latent productive type, $j \in \{1, 2\}$. The effect of size class depends only on a single parameter, ϕ , and we allow the effect of the latent type to be specified fairly generally. Since (s, j) is two-dimensional when n is one-dimensional, and because production must increase in n to apply Shimer's results, our parameterization must sort output, $x_{(m,i),(n,j)}$, in increasing order at every stage of our estimation procedure along the (s, j) dimension.

We estimate several specifications of the production technology that imply different relationships between the supply and demand for skill on the one hand, and the matching of skills to productive characteristics on the other. The LCA technology is additively separable in firm and worker types, and therefore admits comparative advantage for low-ability workers in high-productivity jobs. The NCA technology is multiplicative, and therefore admits no comparative advantage. For all of these parameterizations, we specify that the worker's

skill type, m , and the employer's latent productive type, j , combine thus:

$$\begin{aligned}\bar{x}_{1,1} &= \chi_0 \\ \bar{x}_{2,1} &= \chi_0 + \chi_1 \\ \bar{x}_{1,2} &= \chi_0 + \chi_2 \\ \bar{x}_{2,2} &= \chi_0 + \chi_1 + \chi_2\end{aligned}$$

where $\bar{x}_{m,n}$ is the mean over $(m, i), (n, j)$. To assist identification, we assume that the productive capital of the employer scales linearly with the log of firm size, σ_s . Our empirical strategy is based on estimates of two different specifications for the production technology that embody different assumptions about the nature of comparative advantage in each sector:

- The LCA technology parameterizes output as $f_{LCA}(m, s, j) = \bar{x}_{m,j} + \phi \log(\sigma_s)$
- The NCA technology parameterizes output as $f_{NCA}(m, s, j) = \exp(f_{LCA}(m, s, j))$

χ_0 is fixed exogenously for identification purposes to be 0.1. Since the production technology is required to be increasing in the productive types of both agents, we restrict $\chi_1 > 0$ and $\chi_2 > 0$. In addition to these specifications, we also estimate a quasi-nested model by considering all production functions that can be obtained as a convex combination of the two basic LCA and NCA technologies. That is, we consider a technology of the form:

$$f_{mix}(m, s, j) = \xi f_{LCA}(m, s, j) + (1 - \xi) f_{NCA}(m, s, j).$$

4.2.2 Computation

We fit the parameter $\zeta = (\alpha, \chi_1, \chi_2, \phi, \delta_0, \delta_1, \gamma_0, \gamma_1, \xi)$ to the observed data by the generalized method of moments. While the estimation strategy is standard conceptually, it is complicated by the fact that there is not a general closed-form solution for the coordination

friction equilibrium given an arbitrary set of model primitives. Thus, there is no general analytic closed form for the moment equations, nor for their first- and second-derivatives. We instead use a Monte Carlo approach, simulated annealing, to find the parameter estimate $\hat{\zeta}_A$ that minimizes the generalized distance between the theoretical and empirical moments. This procedure requires sampling the parameter space many times, solving the model at each draw to compute the theoretical moments and thereby the distance from the empirical moments. Solving the coordination friction equilibrium for an arbitrary set of parameters is accomplished by first solving the planner’s problem numerically for the equilibrium queues, and then calculating the wages from the corresponding decentralized economy. The details of these computational methods are outlined below, after first discussing the general setup for the GMM estimation.

Model estimation is complicated by the absence of a closed form expression for the model equilibrium in terms of the parameters. It is not possible to explicitly compute the partial derivatives of the moment equations with respect to the model parameters. We therefore turn to simulation methods to find the GMM estimator. The moment estimator, $\hat{\zeta}_A$ minimizes

$$\hat{\zeta}_A = \arg \min Q_A(\zeta) = \arg \min (\bar{m}(D) - f(\zeta))^T A (\bar{m}(D) - f(\zeta))$$

Simulated annealing is an optimization technique that can be thought of as a variant of greedy search that allows for ‘uphill’ moves. The uphill moves prevent the algorithm from getting stuck in a local minimum. The probability of such a move decreases over the course of the algorithm. During early iterations, the algorithm can easily bounce out of local minima, leading to a more complete exploration of the parameter space.

We use the Adaptive Simulated Annealing (ASA) algorithm described in Ingber (1993).

⁹ Simulated annealing is known to be effective on other estimation problems (Goffe et al.

⁹To implement ASA in MATLAB, we use the **ASAMIN** routine maintained by Shinichi Sakata, as described in Sakata and White (2001).

1994), and the ASA variant of simulated annealing has been successfully applied in a few economic settings. In spite of theoretical assurances of convergence tuning the algorithm to the problem is essential.¹⁰

At each iteration in the algorithm, we must solve the decentralized equilibrium of the coordination friction economy. We exploit the theoretical result from Shimer’s paper that the equilibrium queue lengths are identical in the decentralized economy and in a centralized economy with a planner that faces coordination constraints equivalent to those faced by the workers and employers. At each iteration in the minimization, it is sufficient to solve the planner’s problem by choosing the vector of applicant queues that maximizes the economy’s expected output. Again, because there is no closed form for the planner’s problem, we solve it by numerical methods. This is a constrained nonlinear optimization in 40 decision variables. We have used the built-in Matlab function `fmincon` as well as `KNITRO` to solve this problem at each iteration with comparable results.

Once the simulated annealing procedure converges to an estimate, $\hat{\zeta}_A$, we use standard results to estimate its covariance matrix as:

$$(F^T AF)^{-1} F^T A S A F (F^T AF)^{-1}$$

where $F = F(\zeta) = \frac{\partial f(\zeta)}{\partial \zeta}$. $F(\hat{\zeta}_A)$ is computed using automatic differentiation at the parameter estimate. We compute the sample variance of \bar{m} , denoted S , using a bootstrap procedure. From the full data, we compute \bar{m} on 1,000 bootstrap samples and compute the empirical covariance matrix for the elements of m . We continue to impose the restriction that there is no correlation between the sample moments across size classes, except for the estimates of $\bar{m}^S(D)$.

¹⁰To ensure that our results are not an artifact of the ASA algorithm, we have also estimated the model using a more standard non-linear solver from the `KNITRO` package

4.2.3 Identification

The parametric restrictions on the relationship between employer size and the number of vacancies and the production technology yield an order condition necessary for identification. It remains to be shown that model satisfies the sufficient condition for identification:

$$E(\bar{m}(D)) - f(\zeta) = 0 \Leftrightarrow \zeta = \zeta_0$$

where ζ_0 is the true parameter of the data generating process. We take note of a theoretical result which states that as long as $F(\zeta) = \frac{\partial f(\zeta)}{\partial \zeta}$ has constant rank in a neighborhood of ζ_0 and full-column rank at ζ_0 , then ζ_0 is locally identified (Ruud 2000). Noting that the smallest singular value of $F(\zeta_0)$ is the 2-norm distance of F from the space of rank-deficient matrices (Golub and Loan 1996), it is sufficient to check this value is non-zero, since this combined with the continuous differentiability of F in ζ ensures that there is a neighborhood of ζ_0 for which F always has full column rank.

4.3 Key Assumptions

As discussed in the introduction to this section, our method relies on our ability to map the AKM decomposition, which relies on panel data with a lot of mobility, to a static model of labor market matching. Shimer (2005) argues that “...the model must be extended to a dynamic framework if it is to be taken quantitatively seriously.” Our approach demonstrates that this need not be the case. If the static model provides a good description of the distribution from which observed matches are sampled in the panel data, then our estimation strategy is valid. This assumption is not dissimilar from approaches that assume panel data are generated from the steady-state of a given dynamic equilibrium model as is common in the literature (for example, Postel-Vinay and Robin (2002)). Moreover, Shimer’s coordination

friction model was motivated, in part, by the stylized facts that emerged from the AKM decomposition. As it turns out, a dynamic model is not required to complete the connection between the stylized facts, the model, and the data. This is not to say that there is no value in pushing the coordination friction model, both theoretically, and empirically, but it is not necessary, either for the validity or significance of our results.

5 Results

We have obtained results from estimating the model under the LCA technology, NCA technology, and a nested technology in which the output from every match is a convex combination of the LCA and NCA technologies for that match given the same parameter values. In the nested model, the parameter ξ is the weight on the LCA technology. Table 1 shows estimates obtained from estimating the LCA, NCA, and nested models on data for the Professional, Scientific and Technical Services sector and the Manufacturing sector from the LEHD sample. We present results only for the general specification that mixes the two types of production technology. The estimates of the LCA and NCA technologies are equivalent to restrictions on the value of ξ . Where appropriate, we comment on the validity of those assumed restrictions.

5.1 The Professional, Scientific and Technical Services Sector

5.1.1 Model Primitives

The estimated parameters reveal the type of workers, the type of jobs, and the production technology in the sector. We plot the results of simulating the coordination friction equilibrium under the estimated parameters in Figures 2 and 3. We estimate that 86 percent of the workforce is of high ability, a reasonable result given the nature of the sector. These workers

sort themselves into jobs across deciles that are distinguished in terms of productivity. Figure 2a shows the split between high and low productivity vacancies across firm size deciles. We find that the concentration of high ability workers is not matched by the distribution of vacancies. The majority of vacancies in this sector are of low productivity.

The estimated production technology is illustrated in Figure 2b. Output is measured in implied units that have been scaled by our assumption that $\chi_0 = 0.1$. There is very little productive heterogeneity by employer size. However, there is a productive advantage associated with the latent employer type, but most of the output differentials within this sector are associated with matches of high ability workers into high productivity employers. Output in those matches is roughly four times greater than the output in other matches. Combining the estimated production technology with the estimated distribution of vacancies, we note that this sector is composed of a large low-productivity segment of jobs in which it does not matter much what kind of worker is employed. These are complemented by a small high productivity segment in which there is a substantial advantage to employing high ability workers.

5.1.2 Characteristics of the Estimated Equilibrium

The estimated equilibrium wages and job assignments in the Professional Services sector reflect a small degree of positive assortative matching of workers to jobs. The equilibrium wages, illustrated in Figure 3a, reflect the peculiarities of the coordination friction equilibrium. Note that even though there is a very high production advantage to high ability workers in high ability jobs, the low ability workers actually have a higher equilibrium wage offer from high ability employers. In terms of the model, this occurs because there are so few high productivity jobs and so many high ability workers that competition limits the wages that high ability workers receive. At the same time, the high productivity employers limit their risk of vacancy by making sure even low ability workers are willing to apply for

those jobs. The high wage offers received by low ability workers is compensation for the unemployment risk they face in applying for those jobs. The expected income from applying for such a job has to balance the expected income from applying for a low productivity job, which can almost always be obtained due to the large number of low productivity vacancies relative to the number of job-seekers.

The positive assortative matching result follows from the fact that the estimated production technology is primarily one with no comparative advantage. How the assortativity of matching manifests itself in this sector is revealed in Figures 3d and 3e. The first figure shows the probability that a low ability worker applies to high productivity and low productivity employers. The second shows the same probabilities for high ability workers. For both ability types, the application probabilities sum to one. Most of the mass is concentrated on the low productivity jobs since there are many more of them, and the shape is governed by the way vacancies are distributed across the size classes. Note that the probability of application to low productivity jobs is considerably higher for low ability workers. In other words, better workers are more likely to apply for better jobs, and, of course, to receive them conditional on applying for them. These application strategies result in the expected distribution of realized matches shown in Figures 3b and 3c.

The low productivity vacancies are occupied by a mixture of high and low ability workers, as shown in Figure 3b. We also see that there is a large number of unoccupied vacancies in equilibrium. By contrast, high productivity jobs, which are a small fraction of the total number of jobs, are almost entirely occupied in equilibrium by high ability workers. Based on these estimates of the number of occupied vacancies, we estimate that 88 percent of the employed workforce is of high productivity. In addition, fully 85 percent of the filled vacancies are of low productivity.

The estimates give an overall impression of this sector as being split between a large, low productivity segment and a small high productivity segment. In the low productivity

segment, workers' ability differentials are largely irrelevant to production. Furthermore, there are sufficiently many vacancies that the coordination friction is not binding. Workers can be relatively certain of obtaining a job conditional on applying for it. By contrast, the high productivity segment has very few vacancies and pays much higher wages. Thus, many workers queue for, but fail to receive, these jobs. The relatively mild correlation we observe between wage components hides the extent of sorting both because there are relatively few high productivity vacancies and low ability workers. In addition, the low ability workers actually earn higher wages when they are employed in high productivity jobs.

5.2 The Manufacturing Sector

5.2.1 Model Primitives

The estimated parameters in the manufacturing sector are listed in the second column of Table 1. We estimate that 92 percent of the workforce participating in this sector is of low productivity. This sector is characterized by a dominance of low productivity vacancies. Figure 5a shows the distribution of vacancies between high and low productivity jobs across deciles of the employer size distribution. Here though, the estimates suggest that most of the high productivity vacancies are concentrated in smaller firms, rather than being evenly distributed by size. The mixing parameter on the production technology, ξ is essentially equal to 1, providing strong evidence that there is a comparative advantage for low ability workers in high productivity jobs. The estimated output, illustrated in Figure 5b indicates that both worker and employer types have a relevant impact on output.

5.2.2 Characteristics of the Estimated Equilibrium

The estimated equilibrium wages (Figure 6a) show that high ability workers receive the same wage in most jobs that they actually occupy in equilibrium. By contrast, low ability workers

receive a wage premium when they obtain high productivity jobs. The application behaviors of both types of workers reflect the incentives created by these wages and the nature of comparative advantage that is ultimately reflected in the equilibrium assignment of workers to jobs. Low ability workers are more likely to apply to high productivity jobs than are high ability workers. As a result, we see in Figure 6c that a large portion of the high productivity vacancies are filled by low ability workers.

When we consider the estimated equilibrium job matches, we find that 88 percent of the workforce employed in this segment is of low ability, and 98 percent of the occupied job vacancies are of low productivity. Therefore, there is very little variation, especially for larger firms in employer type. Hence, the small correlations between estimated wage components are an artifact of the very small amount of heterogeneity in employer types in the sector. Nevertheless, the model does a rather good job of capturing the empirical correlation between $\hat{\theta}$ and $\hat{\psi}$. Figure 7 shows the estimated and actual correlations. The model captures both the overall location of these correlations as well as some of the variation across deciles of the employer size distribution.

6 Conclusion

We implemented a formal test of Shimer’s model of job assignment with coordination frictions. The full structural model was estimated using the method of moments. The model predicts the equilibrium number of job matches as well as the empirical moments of the joint distribution of estimated person and employer effects from a linear decomposition of log labor earnings. Our estimates of this model for two major sectors of the U.S. economy provide evidence that there is matching in the labor market in terms of comparative advantage, but the empirical relevance of matching is attenuated because there is insufficient heterogeneity of workers and employers for assignment to influence the distribution of earnings.

This paper provides some of the first estimates of a structural job assignment model with transferable surplus using matched employer-employee data. These estimates allow a structural interpretation of the observed correlations between person and employer effects from the log wage decomposition. While we have focused in this paper on estimating the assignment model with coordination friction, our empirical strategy is general. Any model that delivers a wage offer function which depends on both worker and employer type and which predicts the way in which workers match to jobs in equilibrium could be estimated using this strategy.

Estimates from both sectors, Professional Services and Manufacturing, have similar characteristics. The underlying production technology should generate comparative advantage. In the case of Professional Services, there is positive assortative matching, while in Manufacturing the estimated production technology has a comparative advantage for low ability workers in high productivity employers. In both cases, the influence of sorting by comparative advantage is limited because there are very few of the workers who should be sorting into low productivity jobs. Likewise, there is a very limited number of high productivity vacancies. In other words, each sector has a small, high productivity segment that generates unemployment based on queueing.

Bibliography

- Abowd, J. M., Creecy, R. H. and Kramarz, F. (2002). Computing person and firm effects using linked longitudinal employer-employee data, *Technical Report TP-2002-06*, LEHD, U.S. Census Bureau.
- Abowd, J. M., Gittings, K., McKinney, K. L., Stephens, B. E., Vilhuber, L. and Woodcock, S. (2012). Dynamically consistent noise infusion as confidentiality protection, *Proceedings of the Federal Committee on Statistical Methodology*, Federal Committee on Statistical Methodology.
- Abowd, J. M., Haltiwanger, J. C. and Lane, J. I. (2004). Integrated Longitudinal Employee-Employer Data for the United States, *American Economic Review* **94**(2).
- Abowd, J. M., Kramarz, F. and Margolis, D. N. (1999). High wage workers and high wage firms, *Econometrica* **67**(2): 251–333.
- Abowd, J. M., Lengermann, P. and McKinney, K. L. (2003). The measurement of human capital in the U.S. economy, *Technical Report TP-2002-09*, LEHD, U.S. Census Bureau.
- Abowd, J. M., Stephens, B. E., Vilhuber, L., Andersson, F., McKinney, K. L., Roemer, M. and Woodcock, S. (2009). *Producer Dynamics: New Evidence from Micro Data*, Chicago: University of Chicago Press for the National Bureau of Economic Research, chapter The LEHD Infrastructure Files and the Creation of the Quarterly Workforce Indicators, pp. 149–230.
- Becker, G. S. (1973). A theory of marriage: Part I, *Journal of Political Economy* **81**(4): 813–46.
- Bjelland, M. J. (2007). *Empirical Analyses of Job Displacements and Productivity*, PhD thesis, Cornell University.
- Eeckhout, J. and Kircher, P. (2011). Identifying sorting-in theory, *Review of Economic Studies* **78**(3): 872–906.
- Goffe, W. L., Ferrier, G. D. and Rogers, J. (1994). Global optimization of statistical functions with simulated annealing, *Journal of Econometrics* **60**: 65–99.
- Golub, G. and Loan, C. (1996). *Matrix computations*, Johns Hopkins studies in the mathematical sciences, Johns Hopkins University Press.
- Ingber, L. (1993). Adaptive simulated annealing (ASA) global optimization C-code, *Technical report*, Caltech Alumni Association.
- Iranzo, S., Schivardi, F. and Tosetti, E. (2008). Skill dispersion and firm productivity: An analysis with matched employer-employee data, *Journal of Labor Economics* **26**: 247–285.

- le Maire, D. and Scheur, C. (2008). Job sampling and sorting. Working Paper.
- Lentz, R. (2010). Sorting by search intensity, *Journal of Economic Theory* **145**(4): 1436–1452.
- Lopes de Melo, R. (2009). Sorting in the labor market: Theory and measurement. Working Paper.
- Mendes, R., van den Berg, G. J. and Lindeboom, M. (2010). An empirical assessment of assortative matching in the labor market, *Labour Economics* **17**(6): 919–929.
- Mortensen, D. T. (2003). *Wage Dispersion: Why Are Similar Workers Paid Differently*, Zeuthen lecture book series, MIT Press.
- Postel-Vinay, F. and Robin, J.-M. (2002). Wage dispersion with worker and employer heterogeneity, *Econometrica* **70**: 2295–350.
- Roy, A. D. (1951). Some thoughts on the distribution of earnings, *Oxford Economic Papers* **3**: 135–146.
- Ruud, P. A. (2000). *An Introduction to Classical Econometric Theory*, Oxford University Press.
- Sakata, S. and White, H. (2001). S-estimation of nonlinear regression models with dependent and heterogeneous observations, *Journal of Econometrics* **103**: 5–72.
- Sattinger, M. (1993). Assignment models of the distribution of earnings, *Journal of Economic Literature* **31**: 831–880.
- Shimer, R. (2005). The assignment of workers to jobs in an economy with coordination frictions, *Journal of Political Economy* **113**: 996.
- Stevens, D. W. (2002). Employment that is not covered by state unemployment, *Technical Report TP-2002-14*, LEHD, U.S. Census Bureau.
- Woodcock, S. (2008). Wage differentials in the presence of unobserved worker, firm, and match heterogeneity, *Labour Economics* **15**: 772–794.

Appendices

A An Economy with Heterogeneous Workers and Firms

In this Appendix we derive worker and firm effects in an economy with only two types of workers and a continuum of firms differing by size and production technology. Notation in this appendix differs slightly from the notation in the main paper in order to accommodate the more general specification.

A.1 The Economy

Let the economy consist of two types of workers. This is not as limiting as it sounds, for it may only represent heterogeneity inside a particular class of workers since we consider the effects net of observables. The economy is divided in a finite number \bar{k} of industries. The workers are hired by a continuum of firms, indexed by the share x of type 1 workers each firm hires, the size t of the firm, and the industry k . The share x is unobserved by the econometrician. In industry k , there is a mass (density) $\mu_k(x, t)$ of firms of type (k, x, t) . We omit from what follows the k index except when it is necessary.

The economy as a whole is represented by the space $\Omega = \{\Omega_1, \dots, \Omega_{\bar{k}}\}$ with $\Omega_k = [0, 1] \times \mathbb{R}^+$, the space defining each industry. There is a quantity (or number) $J = \int_{\Omega} d\mu$ of firms, $J_k = \int_{\Omega_k} d\mu_k$ in each industry. Firm (x, t) hires xt of type 1 workers, and $(1 - x)t$ of type 2. We write $\int_{\Omega} xt d\mu(x, t) = M$ and $\int_{\Omega} (1 - x)t d\mu(x, t) = N$ so that there are $M + N$ workers in the economy. Type 1 workers earn $w(x, t, k)$ in firm (x, t) in industry k while type 2 earn $u(x, t, k)$; the wage only depends on the industry, the share of type 1 workers, and the size of the firm.

A.2 The Economy as Estimated by Person and Firm Effects

As in Abowd et al. (1999) (hereafter, AKM99), decompose the individual wage, $\ln w_{i,j(i,t),t}$, after conditioning for observables, in each period as the sum of an individual effect, θ_i , a firm effect, $\psi_{j(i,t)}$, and a residual, $\varepsilon_{i,j(i),t}$

$$\ln w_{i,j(i,t),t} = \theta_i + \psi_{j(i,t)} + \varepsilon_{i,j(i),t}$$

We consider here only the limiting (and purely theoretical) case where the number of time periods goes to infinity and the workers visit ergodically all firms and all industries. This frees us from the identification problem discussed in detail in the above paper. Moreover, the effects estimated in AKM99 would converge in probability to the theoretical values we show below.

In this case, even though the type of the worker is not directly known by the econometrician, observing the workers jump from job to job provides enough information to be able to reduce the above ordinary least squares (OLS) problem to the following minimization

problem. Find W , U (worker effects) and $\psi(k, x, t)$ (firm effect in industry k for a firm of size t hiring a proportion x of type 1 workers) that solve

$$\min_{W, U, \psi} \int_{\Omega} [xt(w(x, t) - W - \psi(x, t))^2 + (1 - x)t(u(x, t) - U - \psi(x, t))^2] d\mu(x, t). \quad (12)$$

This problem is exactly equivalent to minimizing the sum of squared residuals in ordinary least squares. The solutions to this problem are unique up to a constant. We show that:

Theorem 2 Under integrability assumptions on w and u , the solutions to the OLS program (12) are of the type

$$W = \frac{\int x(1 - x)tw(x, t)}{\int x(1 - x)t} + \Lambda, \quad (13)$$

$$U = \frac{\int x(1 - x)tu(x, t)}{\int x(1 - x)t} + \Lambda, \quad (14)$$

$$\psi(x, t) = x(w(x, t) - W) + (1 - x)(u(x, t) - U) - \Lambda \quad (15)$$

where the integrals are all taken on Ω with respect to the measure μ , and Λ is an arbitrary constant (we will set $\Lambda = 0$ in the rest of the paper).

Proof. Set $U = 0$ for the time being, as the solutions are only identified up to a constant. Let us write $J(W, \psi) = \int_{\Omega} \rho(W, \psi)(x, t) d\mu(x, t)$. J is defined on $\mathbb{R} \times L^2(\Omega, \mu)$, and may take the value $+\infty$. It is clear that ρ is convex in its two arguments (as sum of two squares); hence J is also convex, and can be show to be strictly convex if $w(x, t) - u(x, t)$ is not a constant. Hence it has a unique minimum under this assumption, characterized by the first order constraints

$$\int_{\Omega} xt(w(x, t) - W - \psi(x, t)) d\mu(x, t) = 0 \quad (16)$$

$$\int_{\Omega} xt(w(x, t) - W - \psi(x, t)) + (1 - x)t(u(x, t) - \psi(x, t)) = 0 \quad (17)$$

Hence $W = \int xt(w - \psi) d\mu / \int xt dB$ and $\psi = x(w - W) + (1 - x)u$. Replacing ψ by this last value in equation 16 and simplifying with the fact that $\int x^2t = M - \int x(1 - x)t$ and $\int x^2tw = \int xtw - \int x(1 - x)w$ proves the theorem. ■

Remark 3 (Some special cases) If all firms hire the same number of both types of workers, the distribution μ_k for all k is $\delta_{1/2, 2}$, a Dirac distribution of firms of equal size and same number of type 1 workers. This leads to $W = \sum_k w_k / \bar{k}$, $U = \sum_k u_k / \bar{k}$, $\psi_k = \frac{1}{2}(w_k - \bar{w} + u_k - \bar{u})$, where w_k is the wage in industry k .

The case where the distribution μ has all mass in t at $t = 2$ and $w(x, t) = u(x, t) + \delta$ leads to $W = U + \delta$ and $\psi(x, t) = u(x) - U$.

We define the firm-employee effects covariances $\text{cov}(\hat{\theta}, \hat{\psi})$ as

$$\frac{\int (xtW + (1-x)tU)\psi(x, t) d\mu}{\int t d\mu} - \frac{\int (xtW + (1-x)tU) d\mu}{\int t d\mu} \frac{\int \psi(x, t) d\mu}{\int t d\mu} \quad (18)$$

For total covariance the integrals are taken over Ω , for within-sector covariance over Ω_k for all k , and the between-sector covariance is the discrete covariance between the means (i.e. integrals) on the Ω_k . Total covariance is as usual the sum of between-sector and within-sector covariances.

B Data Appendix

We use data from the Longitudinal Employer-Household Dynamics Program (LEHD), universe data for twenty-eight large American states with information from 1990-2004. See Abowd et al. (2004) and Abowd et al. (2009) for more detailed discussions of the underlying data. Our analysis sample is restricted to individuals aged 18-70, employed full-time at their dominant employer. The following sections describe the underlying microdata sources and how we construct the analysis variables and impose sample restrictions.

B.1 Unemployment Insurance

The individual data were derived from the universe of unemployment insurance (UI) quarterly wage records from twenty-eight states.¹¹ The BLS Handbook of Methods (1997) describes UI coverage as “broad and basically comparable from state to state,” and claims over 96 percent of total wage and salary civilian nonfarm jobs were covered in 1994. The Federal Unemployment Tax Act (FUTA), first enacted in 1938, lays the ground rules for the kinds of employment which must be covered in state unemployment insurance laws. While technically mandating coverage of all employers with one or more employees in a calendar year, FUTA allows for numerous exceptions to covered employment (Stevens 2002). These include workers at small agricultural co-operatives, employees of the federal government, and certain employees of state governments, most notably elected officials, members of the judiciary, and emergency workers. According to the Handbook, UI wage records measure “gross wages and salaries, bonuses, stock options, tips, and other gratuities, and the value of meals and lodging, where supplied.” They do not include OASDI, health insurance, workers compensation, unemployment insurance, and private pension and welfare funds. Individuals are uniquely identified and followed for all quarters in which their employers had reporting requirements in the UI system. Thus, cross-state mobility can be observed for individuals moving between any of states for which we have data. Although coverage dates vary, all states provide between seven and thirteen years of data. In our analysis, by combining them

¹¹We cannot disclose which states are included in the analysis as a condition imposed by the Census Bureau on the public use data. We can, however, release the moment matrices used in the structural estimation, upon request, as they have been designated as public-use data.

into a single “pooled” file, we have information on approximately 190 million workers, 10 million employers, and 660 million annual earnings records.

B.2 Creation of Variables

Using Census Bureau and other LEHD data bases, sex, race, and date of birth are combined with the individual earnings data using an exact match to administrative data. Upon each individual’s first appearance in the data, labor force experience as potential labor force experience ($\text{age} - \text{education} - 6$) is calculated. In subsequent periods, experience is measured as the sum of observed experience and initial (potential) experience. The UI wage records connect individuals to every employer from which they received wages in any quarter of a given calendar year. Therefore, individual employment histories are constructed using the same personal identifier used in the individual data. Employers are identified by their state unemployment insurance account number (SEIN). While large employers undoubtedly operate in multiple states, their SEINs are unfortunately state specific, meaning they cannot be connected. In addition, while we match workers to their employers, it is not possible to connect those employed in firms with multiple establishments to specific places of work. This problem is not overly pervasive, as over 70% of employment occurs in firms with only a single establishment.

B.3 Earnings

For every year an individual appears in a state, a “dominant” employer — the employer for whom the sum of quarterly earnings is the highest—is identified in order to better approximate the individual’s full-time, full-year annual wage rate using the following steps. First, define full quarter employment in quarter t as having an employment history with positive earnings for quarters $t - 1, t$, and $t + 1$. Continuous employment during quarter t means having an employment history with positive earnings for either $t - 1$ and t or t and $t + 1$. Employment spells that are neither full quarter nor continuous are designated discontinuous. If the individual was full quarter employed for at least one quarter at the dominant employer, the annualized wage is computed as 4 times average full quarter earnings at that employer (total full quarter earnings divided by the number of full quarters worked). This accounts for 84% of the person-year-state observations in our eventual analysis sample. Otherwise, if the individual was continuously employed for at least one quarter at the dominant employer, the annualized wage is average earnings in all continuous quarters of employment at the dominant employer multiplied by 8 (i.e., 4 quarters divided by an expected employment duration during the continuous quarters of 0.5). This accounts for 11% of all observations. For the remaining 5%, annualized wages are average earnings in each quarter multiplied by 12 (i.e., 4 quarters divided by an expected employment duration during discontinuous quarters of 0.33).

B.4 Annual Hours of Work and Full-time Status

We restrict our sample to individuals who worked full-time for their dominant employer. Full-time status is taken to mean that the individual worked at least 35 hours per week. The average number of hours per week for a worker is given by dividing the number of hours worked over the year with the dominant employer by the number of weeks worked. The number of weeks worked is approximated based on the observed number of full, continuous, and discontinuous quarters with one's dominant employer. Specifically, we compute weeks worked as 13 times the number of full quarters + 6.5 times the number of continuous quarters + 4.33 times the number of discontinuous quarters. Annual hours of work are observed for a subset of individuals in the sample via a database link to other LEHD and Census sources. For those missing annual hours of work, a value is statistically assigned using a Bayesian multiple imputation procedure. Bjelland (2007) provides complete details of the imputation of annual hours of work.

B.5 Human Capital Estimates


To the dominant job frame, we attach the variance components of earnings estimated from the Abowd et al. (1999) human capital decomposition. The estimation of those variance components in the LEHD data is the subject of another paper, Abowd et al. (2002), to which we refer the interested reader.

C Tables and Figures

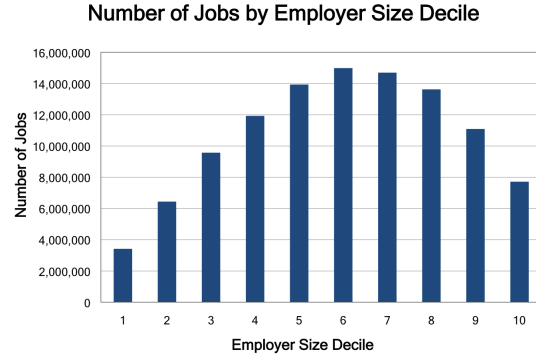
Table 1: Estimates of the Model Parameters

Parameter	Prof. Svcs.	Mfg.
α	0.139 (0.0001)	0.924 (0.0003)
χ_1	2.133 (0.0007)	0.767 (0.0009)
χ_2	3.829 (0.0116)	3.106 (0.0030)
ϕ	0.027 (0.0004)	0.0282 (0.0001)
δ_0	2.322 (0.0063)	0.358 (0.0047)
δ_1	0.022 (0.0009)	1.034 (0.0021)
γ_0	1.681 (0.0507)	-0.276 (0.0321)
γ_1	-2.216 (0.0376)	-1.616 (0.0102)
ξ	0.105 (0.0048)	0.999 (0.0000)
Residual	0.0900	0.0214


GMM parameter estimates for the nested model that mixes the two technologies according to the parameter ξ . See text for details. Standard errors in parentheses.




(a) Prof. Svcs.




(b) Manufacturing




(c) Prof. Svcs.



(d) Manufacturing




(e) Prof. Svcs.

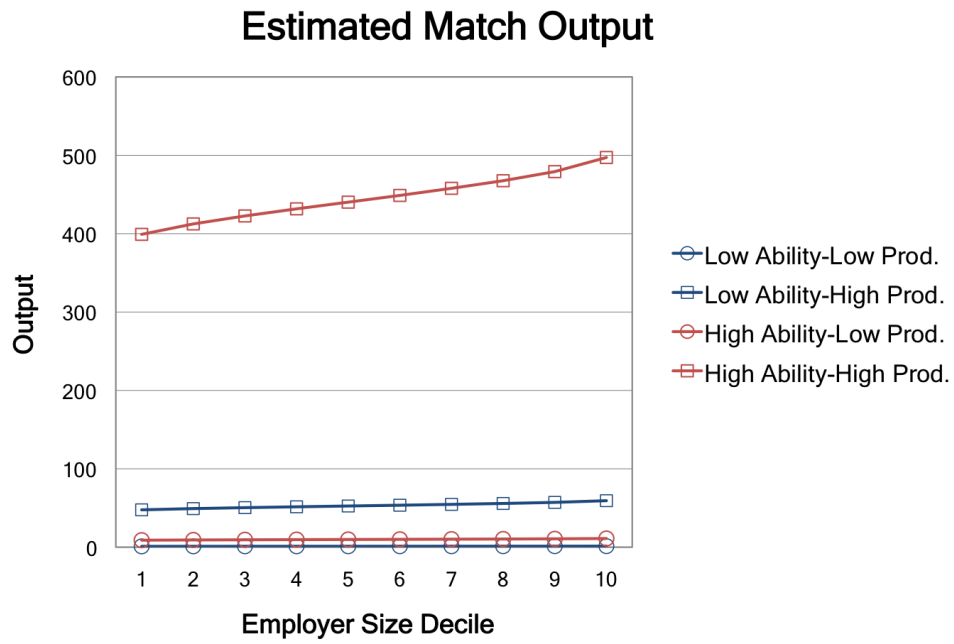


(f) Manufacturing

Figure 1: Data used in Estimation: Professional Services (NAICS 54) and Manufacturing (NAICS 31–33)



(a) Vacancies




(b) Production

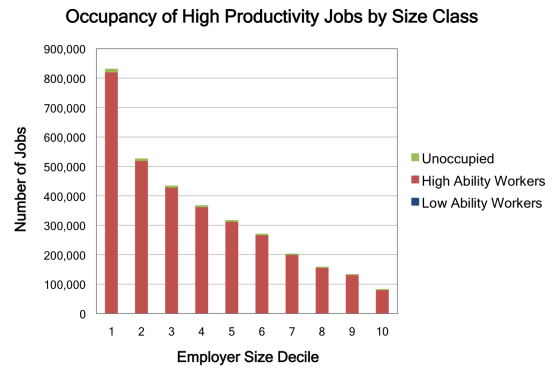
Figure 2: Estimated Model Primitives: Professional Services Sector




(a) Wages



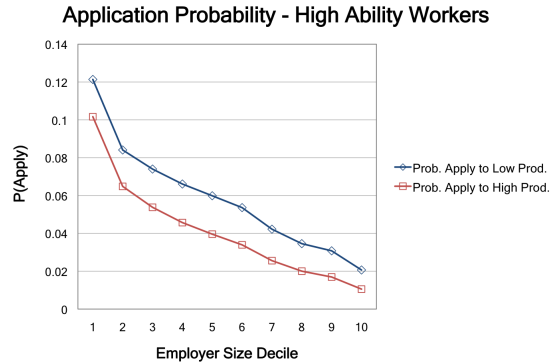
(b) Low Prod. Vacs.



(c) High Prod. Vacs.




(d) App. Prob. Low Ability



(e) App. Prob. High Ability


Figure 3: Estimated Equilibrium: Professional Services Sector

Variance of Person Effects - Actual and Estimated



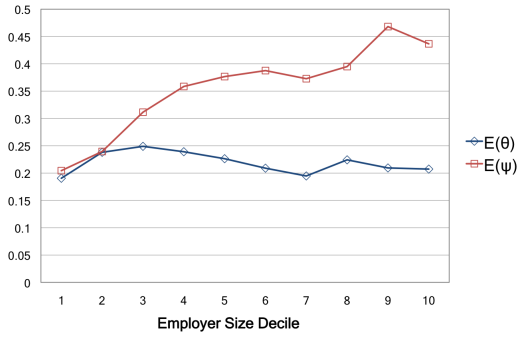
(a) $\text{Var}(\theta)$

Variance of Firm Effects - Actual and Estimated




(b) $\text{Var}(\psi)$

Empirical Average Person and Firm Effects




(c) Empirical $E(\theta), E(\psi)$

Estimated Average Person and Firm Effect




(d) Estimated $E(\theta), E(\psi)$

Correlation of Person and Firm Effects

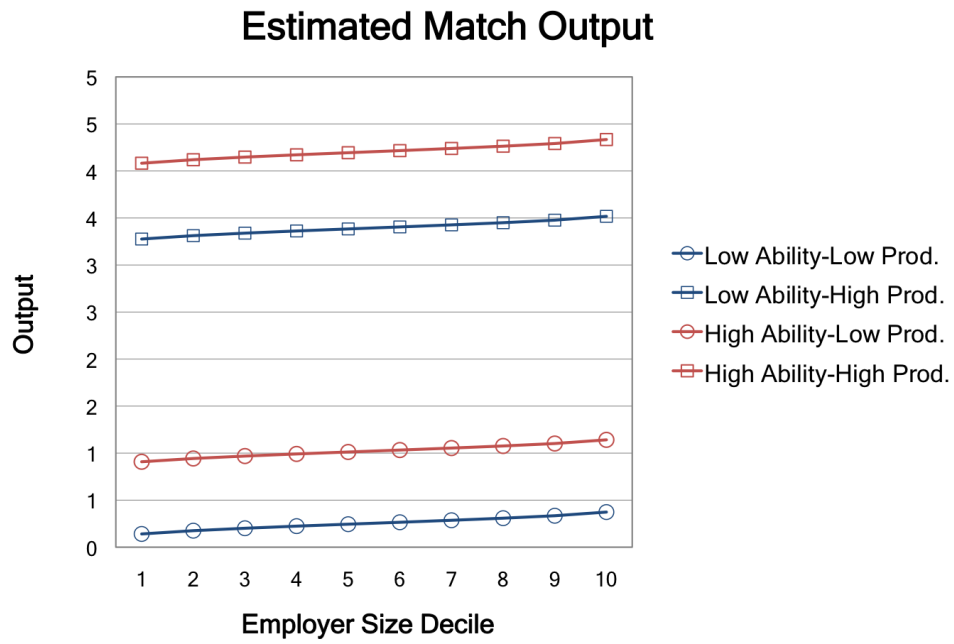


(e) $\text{Corr}(\theta, \psi)$

Figure 4: Data and Model Fit: Professional Services Sector




(a) Vacancies




(b) Production


Figure 5: Estimated Model Primitives: Manufacturing Sector




(a) Wages




(b) Low Prod. Vacs.



(c) High Prod. Vacs.




(d) App. Prob. Low Ability



(e) App. Prob. High Ability


Figure 6: Estimated Equilibrium: Manufacturing Sector

Variance of Person Effects - Actual and Estimated



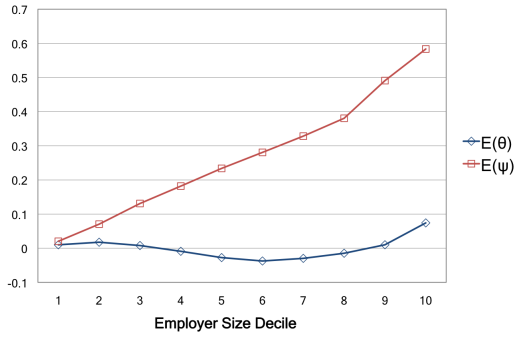
(a) $\text{Var}(\theta)$

Variance of Firm Effects - Actual and Estimated




(b) $\text{Var}(\psi)$

Empirical Average Person and Firm Effects




(c) Empirical $E(\theta), E(\psi)$

Estimated Average Person and Firm Effect



(d) Estimated $E(\theta), E(\psi)$

Correlation of Person and Firm Effects



(e) $\text{Corr}(\theta, \psi)$

Figure 7: Data and Model Fit: Manufacturing Sector